

0607_boolgan

September 27, 2020

```
[ ]: # Don't run again  
  
# !unzip 'thecarconnectionpicturedataset.zip' -d 'dataset/'
```

```
[ ]: from __future__ import print_function  
#%matplotlib inline  
import argparse  
import os  
import random  
import torch  
import torch.nn as nn  
import torch.nn.parallel  
import torch.backends.cudnn as cudnn  
import torch.optim as optim  
import torch.utils.data  
import torchvision.datasets as dset  
import torchvision.transforms as transforms  
import torchvision.utils as vutils  
import numpy as np  
import matplotlib.pyplot as plt  
import matplotlib.animation as animation  
from IPython.display import HTML  
from PIL import Image  
from fid_score import calculate_fid_given_paths  
import time  
  
# Set random seed for reproducibility  
manualSeed = 999  
#manualSeed = random.randint(1, 10000) # use if you want new results  
print("Random Seed: ", manualSeed)  
random.seed(manualSeed)  
torch.manual_seed(manualSeed)
```

Random Seed: 999

```
[ ]: <torch._C.Generator at 0x7f62c586f230>
```

```
[ ]: dataroot = "dataset/"

# Number of workers for dataloader
workers = 2

# Batch size during training
batch_size = 128

# Spatial size of training images. All images will be resized to this
# size using a transformer.
image_size = 64

# Number of channels in the training images. For color images this is 3
nc = 3

# Size of z latent vector (i.e. size of generator input)
nz = 100

# Size of feature maps in generator
ngf = 64

# Size of feature maps in discriminator
ndf = 64

# Number of training epochs
num_epochs = 50

# Learning rate for optimizers
lr = 0.00075

# Beta1 hyperparam for Adam optimizers
beta1 = 0.5

# Number of GPUs available. Use 0 for CPU mode.
ngpu = 1
```

```
[ ]: # We can use an image folder dataset the way we have it setup.
# Create the dataset
dataset = dset.ImageFolder(root=dataroot,
                           transform=transforms.Compose([
                               transforms.Resize(image_size),
                               transforms.CenterCrop(image_size),
                               transforms.ToTensor(),
                               transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
                           ]))
# Create the dataloader
```

```
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
                                         shuffle=True, num_workers=workers)

# Decide which device we want to run on
device = torch.device("cuda:0" if (torch.cuda.is_available() and ngpu > 0) else "cpu")

# Plot some training images
real_batch = next(iter(dataloader))
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")
print(real_batch[0].to(device)[:64].shape)
plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:64],  
                         padding=2, normalize=True).cpu(),(1,2,0)))
```

```
torch.Size([64, 3, 64, 64])
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f62c17dde80>
```

Training Images



```
[ ]: for i in range(64):
    vutils.save_image(real_batch[0][i], \
                      os.path.join('sample_real_images_0607',"foo"+str(i)+".
→jpeg"), \
                      normalize=True)
```

```
[ ]: print(torch.cuda.is_available())
print(torch.cuda.device_count())
```

True
1

```
[ ]: # custom weights initialization called on netG and netD
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)
```

```
[ ]: # Generator Code
```

```
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d(nz, ngf * 16, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 16),
            nn.LeakyReLU(True),
            # state size. (ngf*16) x 4 x 4
            nn.ConvTranspose2d(ngf * 16, ngf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 8 x 8
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 16 x 16
            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. (ngf*2) x 32 x 32
            nn.ConvTranspose2d(ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),

            # state size. (ngf) x 64 x 64
            nn.ConvTranspose2d(ngf, nc, 4, 2, 1, bias=False),
            nn.BatchNorm2d(nc),
            nn.ReLU(True),

            # state size. (nc) x 128 x 128
            nn.Conv2d(nc, nc * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(nc * 2),
            nn.ReLU(True),
```

```

        # state size. (nc * 2) x 64 x 64
        nn.Conv2d(nc * 2, nc, 1, 1, 0, bias=False),
        nn.Tanh()
    )

    def forward(self, input):
        return self.main(input)

```

```

[ ]: # Create the generator
netG = Generator(ngpu).to(device)

# Handle multi-gpu if desired
if (device.type == 'cuda') and (ngpu > 1):
    netG = nn.DataParallel(netG, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, stdev=0.2.
netG.apply(weights_init)

# Print the model
print(netG)

```

```

Generator(
    main): Sequential(
        (0): ConvTranspose2d(100, 1024, kernel_size=(4, 4), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): LeakyReLU(negative_slope=True)
        (3): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
        (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
        (7): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (8): ReLU(inplace=True)
        (9): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
        (10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (11): ReLU(inplace=True)
        (12): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
        (13): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
(14): ReLU(inplace=True)
(15): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1), bias=False)
(16): BatchNorm2d(3, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(17): ReLU(inplace=True)
(18): Conv2d(3, 6, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
(19): BatchNorm2d(6, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(20): ReLU(inplace=True)
(21): Conv2d(6, 3, kernel_size=(1, 1), stride=(1, 1), bias=False)
(22): Tanh()
)
)

```

```

[ ]: class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 8 x 8
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Dropout(0.2),
        )

    def forward(self, input):
        return self.main(input)

```

```
[ ]: # Create the Discriminator
netD = Discriminator(ngpu).to(device)

# Handle multi-gpu if desired
if (device.type == 'cuda') and (ngpu > 1):
    netD = nn.DataParallel(netD, list(range(ngpu)))

# Apply the weights_init function to randomly initialize all weights
# to mean=0, std=0.2.
netD.apply(weights_init)

# Print the model
print(netD)
```

```
Discriminator(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Dropout(p=0.2, inplace=False)
  )
)
```

```
[ ]: # Initialize BCELoss function
criterion = nn.BCELoss()

# Create batch of latent vectors that we will use to visualize
# the progression of the generator
fixed_noise = torch.randn(64, nz, 1, 1, device=device)

# Establish convention for real and fake labels during training
```

```

real_label = 1
fake_label = 0

# Setup Adam optimizers for both G and D
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

```

[]: # Training Loop

```

# Lists to keep track of progress
img_list = []
G_losses = []
D_losses = []
fid_values = []
iters = 0

print("Starting Training Loop...")
# For each epoch
for epoch in range(num_epochs):
    # For each batch in the dataloader
    for i, data in enumerate(dataloader, 0):

        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        ## Train with all-real batch
        netD.zero_grad()
        # Format batch
        real_cpu = data[0].to(device)
        b_size = real_cpu.size(0)
        label = torch.full((b_size,), real_label, device=device)
        # Forward pass real batch through D
        output = netD(real_cpu).view(-1)
        D_x = output.mean()

        ## Train with all-fake batch
        # Generate batch of latent vectors
        noise = torch.randn(b_size, nz, 1, 1, device=device)
        # Generate fake image batch with G
        fake = netG(noise)
        label.fill_(fake_label)
        # Classify all fake batch with D
        output = netD(fake.detach()).view(-1)
        D_G_z1 = output.mean()#.item()
        errD = D_G_z1 - D_x
        errD.backward()
        # Update D

```

```

optimizerD.step()
# Wasserstein weight clipping
for p in netD.parameters():
    p.data.clamp_(-0.1, 0.1)

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake
→batch through D
    output = netD(fake).view(-1)
    # Calculate G's loss based on this output
    D_G_z2 = output.mean()
    errG = -D_G_z2
    errG.backward()
    # Update G
    optimizerG.step()

    # Output training stats
    if i % 50 == 0:
        print(' [%d/%d] [%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.
→4f\tD(G(z)): %.4f / %.4f'
            % (epoch, num_epochs, i, len(dataloader),
               errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))

    # Save Losses for plotting later
    G_losses.append(errG.item())
    D_losses.append(errD.item())

    # Check how the generator is doing by saving G's output on fixed_noise
    if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i ==
→len(dataloader)-1)):
        with torch.no_grad():
            fake = netG(fixed_noise).detach().cpu()
            for i in range(len(fake)):
                vutils.save_image(fake[i], \
                    os.path.join('sample_fake_images_0607', "fake"+str(i)+".
→jpeg"), \
                    normalize=True)
            time.sleep(30)
            fid_value = calculate_fid_given_paths(['sample_real_images_0607/
→', 'sample_fake_images_0607/'], 50, False, 2048)
            fid_values.append(fid_value)
            print("FID Value:", fid_value)

```

```

    img_list.append(vutils.make_grid(fake, padding=2, normalize=True))

    iters += 1

```

Starting Training Loop...

[0/50] [0/155]	Loss_D: -71.5861 -29.1921 / -58.2906	Loss_G: 58.2906 D(x): 42.3939	D(G(z)):
[0/50] [50/155]	Loss_D: -91.0381 -43.1498 / -57.1893	Loss_G: 57.1893 D(x): 47.8883	D(G(z)):
[0/50] [100/155]	Loss_D: -85.4512 -50.1064 / -37.8022	Loss_G: 37.8022 D(x): 35.3447	D(G(z)):

0% | 0/2 [00:00<?, ?it/s]

['sample_real_images_0607/', 'sample_fake_images_0607/']
[PosixPath('sample_real_images_0607/foo0.jpeg'),
 PosixPath('sample_real_images_0607/foo1.jpeg'),
 PosixPath('sample_real_images_0607/foo2.jpeg'),
 PosixPath('sample_real_images_0607/foo3.jpeg'),
 PosixPath('sample_real_images_0607/foo4.jpeg'),
 PosixPath('sample_real_images_0607/foo5.jpeg'),
 PosixPath('sample_real_images_0607/foo6.jpeg'),
 PosixPath('sample_real_images_0607/foo7.jpeg'),
 PosixPath('sample_real_images_0607/foo8.jpeg'),
 PosixPath('sample_real_images_0607/foo9.jpeg'),
 PosixPath('sample_real_images_0607/foo10.jpeg'),
 PosixPath('sample_real_images_0607/foo11.jpeg'),
 PosixPath('sample_real_images_0607/foo12.jpeg'),
 PosixPath('sample_real_images_0607/foo13.jpeg'),
 PosixPath('sample_real_images_0607/foo14.jpeg'),
 PosixPath('sample_real_images_0607/foo15.jpeg'),
 PosixPath('sample_real_images_0607/foo16.jpeg'),
 PosixPath('sample_real_images_0607/foo17.jpeg'),
 PosixPath('sample_real_images_0607/foo18.jpeg'),
 PosixPath('sample_real_images_0607/foo19.jpeg'),
 PosixPath('sample_real_images_0607/foo20.jpeg'),
 PosixPath('sample_real_images_0607/foo21.jpeg'),
 PosixPath('sample_real_images_0607/foo22.jpeg'),
 PosixPath('sample_real_images_0607/foo23.jpeg'),
 PosixPath('sample_real_images_0607/foo24.jpeg'),
 PosixPath('sample_real_images_0607/foo25.jpeg'),
 PosixPath('sample_real_images_0607/foo26.jpeg'),
 PosixPath('sample_real_images_0607/foo27.jpeg'),
 PosixPath('sample_real_images_0607/foo28.jpeg'),
 PosixPath('sample_real_images_0607/foo29.jpeg'),
 PosixPath('sample_real_images_0607/foo30.jpeg'),
 PosixPath('sample_real_images_0607/foo31.jpeg'),
 PosixPath('sample_real_images_0607/foo32.jpeg'),

```


-52.4274 / -40.9825
[8/50] [0/155] Loss_D: -68.9025 Loss_G: 55.1833 D(x): 57.6392 D(G(z)):
-11.2632 / -55.1833
[8/50] [50/155] Loss_D: -84.7135 Loss_G: 59.3480 D(x): 61.8538 D(G(z)):
-22.8597 / -59.3480
[8/50] [100/155] Loss_D: -72.1132 Loss_G: 51.5793 D(x): 60.2030 D(G(z)):
-11.9102 / -51.5793
[8/50] [150/155] Loss_D: -73.7943 Loss_G: 27.5061 D(x): 19.4245 D(G(z)):
-54.3699 / -27.5061
[9/50] [0/155] Loss_D: -64.5925 Loss_G: 48.1603 D(x): 49.0829 D(G(z)):
-15.5096 / -48.1603
[9/50] [50/155] Loss_D: -83.8324 Loss_G: 37.4948 D(x): 36.9237 D(G(z)):
-46.9087 / -37.4948
[9/50] [100/155] Loss_D: -69.9681 Loss_G: 58.1766 D(x): 58.7819 D(G(z)):
-11.1862 / -58.1766
[9/50] [150/155] Loss_D: -60.9230 Loss_G: 40.0866 D(x): 10.3662 D(G(z)):
-50.5568 / -40.0866

```

```

[ ]: plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses,label="G")
plt.plot(D_losses,label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
plt.show()


```



```
[ ]: %%capture
fig = plt.figure(figsize=(8,8))
plt.axis("off")
ims = [[plt.imshow(np.transpose(i,(1,2,0)), animated=True)] for i in img_list]
ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)

HTML(ani.to_jshtml())
```

[]: <IPython.core.display.HTML object>




```
[ ]: plt.figure(figsize=(10,5))
plt.title("Frechet Inception Distance During Training of BoolGAN")
```

```

plt.plot(fid_values, marker='o')
plt.xlabel("Iterations")
labelList = range(0, len(fid_values)*500, 500)
plt.xticks(ticks=range(0, len(fid_values)), labels=labelList, rotation=45)
plt.ylabel("FID")
plt.show()

```



```
[ ]: print(fid_values)
```

```
[425.4186359162443, 313.9360441648118, 312.52982531308646, 311.64497167846264,
289.8353060933684, 259.374964764839, 240.29331157594982, 229.3073117959616,
219.0035415450701, 214.94912749317803, 207.9029070554729, 202.0819681655023,
196.3733133334875, 189.92914427573209, 192.65490340057005, 184.8520162324376,
184.41251005806075, 177.73243766373858, 181.0522812208194, 182.22659900846674,
178.5358250389574, 171.12284094889984, 171.0830520343335, 165.9662036438666,
175.17854331246605, 169.64499489759885, 174.91007214849043, 171.98485570346128]
```

```
[ ]: # Grab a batch of real images from the dataloader
real_batch = next(iter(dataloader))

# Plot the real images
plt.figure(figsize=(15,15))
plt.subplot(1,2,1)
plt.axis("off")
plt.title("Real Images")
```


```

plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:, :64],  

                                         padding=5, normalize=True).cpu(), (1, 2, 0)))

# Plot the fake images from the last epoch
plt.subplot(1, 2, 2)
plt.axis("off")
plt.title("Fake Images")
plt.imshow(np.transpose(img_list[-1], (1, 2, 0)))
plt.show()

```



[]: